

# Chronicle:

## An Exploration in Durable Session Memory for Claude Code

Fergus

May 2026

### Abstract

Stock Claude Code is excellent at searching the repository in front of it. Chronicle targets the missing substrate: prior agent sessions. It is a hook-delivered memory layer that retrieves project context from the user's own session JSONLs and renders a `<chronicle-context>` block into the next turn before the model has to know it exists. This write-up presents three layers of evidence, all run on the same harness with Docker isolation, OAuth-only execution authentication, and Chronicle served by a per-tuple sidecar daemon. The first layer is a vague track on open-ended prior-context prompts, scored by concept check rather than the old non-empty-answer floor. The second layer is a behaviour track that includes Headroom as a real competing memory tool. The third layer is the `chronicle_advantage_pack`, a purpose-built three-arm Sonnet pack (baseline, `chronicle-on`, `headroom-on`) designed end to end around the prior-session-dependence assumption, conversation-weighted, with budget dominance as the headline figure. Across the three layers, Chronicle solves more prior-session-dependent tasks per budget than stock Claude Code and Headroom, with lower cost per solved task, lower wall time per solved task, and lower repository-search per solved task. On the advantage pack, Headroom does not lift hard pass over baseline and helps marginally on signal; Chronicle dominates both arms on every per-success metric.

**Keywords:** coding agents; memory; retrieval; evaluation; Claude Code; prior session context; budget dominance.

## 1 The framing this write-up earns

Stock Claude Code can search the repository in front of it. It cannot search what already happened: the design decisions, the rejected alternatives, the project-internal names, the corrections, and the partially completed work that only exist in earlier sessions. Every long-running agentic project hits this gap, and the standard fix, a small curated `MEMORY.md`, does not scale. Chronicle indexes past Claude Code sessions and surfaces the right context to each new prompt, so the agent picks up where prior work left off rather than rediscovering it.

The defensible claim, supported by the three layers of evidence below, is:

On prior-session-dependent tasks, Chronicle solves more tasks per budget than stock Claude Code: lower cost per solved task, lower wall time per solved task, and lower repository-search per solved task. Same-task per-pair cost is modestly positive; the Chronicle win comes from solving more tasks at comparable per-pair spend, not from solving the same task more cheaply.

The headline figure is therefore a budget-dominance curve, not a same-task latency comparison.

## 2 On concurrent work

I started this work on 27 April 2026. On 6 May 2026, ten days into the project, Anthropic announced *Dreaming* as a research preview for Claude Managed Agents. Dreaming reviews an agent’s past sessions, extracts patterns, and curates memories so agents learn over time. The problem framing is the same as Chronicle’s: prior-session context is the missing substrate for agentic systems, and the natural fix is to mine, curate, and retrieve from that history.

I take Anthropic’s announcement as independent confirmation that the problem matters in production agentic systems rather than as the prompt for this work. The two efforts target different surfaces. Dreaming is for Claude Managed Agents, Anthropic’s hosted agent product. Chronicle is for Claude Code, the local CLI agent that writes session JSONLs to `~/.claude/projects/`. The substrates, the latency budgets, and the delivery mechanisms are different: Dreaming runs in Anthropic’s cloud over a managed agent’s history; Chronicle runs locally on the user’s machine, indexing raw JSONL transcripts, and serves retrieval through Claude Code hooks rather than through a managed loop. The convergence on the problem space is useful evidence that the gap is real. The contribution this write-up defends is a local, hook-delivered, JSONL-native implementation specific to Claude Code’s substrate.

## 3 Why Chronicle exists

I started digging into `MEMORY.md` after spending a day switching between two Claude Code accounts, re-explaining the same deployment process in every new chat. What I found was a single markdown file capped at 200 lines or 25 KB, loaded at session start, written to by Claude Code’s Auto Memory feature, and periodically consolidated by AutoDream during idle periods. My current project’s file had one line. The other projects I had been working on were the same story. For me, prior-session knowledge essentially did not exist between turns. The interesting question is whether a per-prompt retrieval layer over the actual session JSONLs can do better than a single curated markdown file, while keeping latency low enough that hooks remain practical.

## 4 How Chronicle works

Chronicle is a per-prompt retrieval layer for Claude Code. It indexes the project’s real session JSONLs (the transcripts Claude Code writes to `~/.claude/projects/<encoded>/<uuid>.jsonl`) and writes a `<chronicle-context>` block into the model’s view of the next turn. The substrate is the user’s own prior sessions; nothing synthetic.

### 4.1 Two retrieval substrates

Chronicle runs two indexes in parallel.

**Substrate A: conversation and action chunks.** The session JSONLs are split into roughly 400-word windows and embedded with BAAI/bge-small-en-v1.5 (33M parameters, 384-d embeddings, ~6 ms encode steady state), then stored in SQLite plus `sqlite-vec`. The spoken-chunk path keeps user/assistant exchange and deliberately excludes `tool_use` markers, `tool_result` content, `thinking` blocks, slash-command echoes, and synthetic tool-result-as-user turns. The chunker also creates behavioral chunks for standard Bash, Edit, Write, and MultiEdit tool inputs and `tool_result` chunks for filtered failures or narrow deploy/publish success evidence. Chunk IDs are content-addressed (`sha256(session || first_turn || content)[:16]`), so unchanged JSONLs re-index for free.

The reasoning: a query about `make_cache_key` should match the discussion of the function, not a code dump that happens to contain it.

**Substrate B: typed-claim memory store.** Markdown files extracted offline from the same JSONLs by Claude Haiku. Each file has metadata at the top (kind, polarity, referent, source\_session, source\_turn) and is indexed twice: FTS5 (BM25) for keyword and sqlite-vec for embedding, combined with Reciprocal Rank Fusion. A graph of replaced memories sits alongside, with links between old and new claims, plus cycle checks at construction time.

## 4.2 Execution loop

A `UserPromptSubmit` event with both substrates enabled currently runs through this shape:

1. chronicle hook reads the payload from stdin (prompt, session\_id, project\_dir, recent\_turns, compaction\_count).
2. A path-triage classifier (SKIP / LIGHT / FULL) inspects prompt shape. SKIP exits silently; LIGHT only adds session-state blocks.
3. Conversation search tries the daemon first. If the daemon is unavailable, the hook falls back to an in-process embedder and store.
4. The query runs through dense, hybrid, or multi-query search depending on config and prompt shape, with min\_top\_cosine, rerank, and MMR controls when enabled.
5. Candidate chunks are re-ranked with polarity, topic overlap, recency, list-answer penalties, stale/replacement signals, and citation-use scores.
6. The typed-claim store is searched alongside the raw chunks; results use kind, polarity, supersession, temporal, and lexical scorers.
7. A routing-weights classifier (sub-millisecond, no LLM) outputs source and memory-kind weights based on prompt patterns.
8. The composer allocates source budgets, removes duplicate overlap, formats blocks, and writes a `<chronicle-context>` block to stdout.
9. Claude Code prepends the block to the model’s view of the turn.

The hook always exits 0; failures inside `run_hook` are caught and logged. The 5-second timeout in `settings.json` is the safety net. Steady-state hot path is tens of milliseconds once the embedder is warm. All three reportable runs in this write-up use the per-tuple sidecar daemon path with TCP transport on macOS.

## 4.3 Hook events used in this write-up

| Event                                  | Purpose                                               | Used in                                                |
|----------------------------------------|-------------------------------------------------------|--------------------------------------------------------|
| <code>UserPromptSubmit</code>          | Front-of-turn injection                               | vague, behaviour, advantage pack<br>behaviour subtrack |
| <code>PreToolUse:Read/Grep/Glob</code> | Mid-turn injection during repository exploration      | advantage pack conversation<br>subtrack                |
| <code>SessionEnd</code>                | Typed-claim extraction over the just-finished session | substrate building                                     |

**Table 1.** Hook events wired in the reportable runs.

## 4.4 Supersession

The supersession graph is built from each memory’s `revises` field. It exposes `is_superseded`, `superseded_by`, `latest_in_chain`, and a bi-temporal `validity_at(t)` projection inspired by Zep / Graphiti. The graph powers a `SUPERSEDED_PENALTY = 0.5` multiplicative downrank during scoring and a `PostToolUse:Read` annotation that catches mid-turn reads of stale memory files even when they

bypass the front-of-turn injection.

## 4.5 Design decisions

A few decisions carry most of the design.

**Hook delivery, not MCP.** MCP would make memory a tool the model chooses to call. Hooks put memory into the prompt path without asking the model to opt in, which matches the [MEMTRACK](#) finding: adding memory tools to LLM agents tends to increase tool-call redundancy and degrade efficiency, with off-the-shelf memory components like Zep and Mem0 not significantly improving performance on long-horizon multi-platform tasks. MCP also adds a round trip on the path where latency matters most.

**LLM-free hot path.** Per-prompt work runs locally: embedding, `sqlite-vec` lookup, heuristic re-rank, triage, and composition. `ANTHROPIC_API_KEY` is only needed for offline typed-claim extraction.

**BGE-small over MiniLM.** A trade-off, not a leaderboard pick. BGE-small worked better on Chronicle’s own session data: short project-specific decisions where the query and source chunk often do not use the same words.

**Separating spoken text from tool output.** File dumps and command output can drown out the discussion around them. The chunker keeps spoken chunks clean, then indexes tool inputs and selected tool results as separate chunk kinds.

## 5 What we measure

Three co-primary endpoints, in this order:

| Endpoint                | Question                                                  |
|-------------------------|-----------------------------------------------------------|
| Hard pass               | Did the task get done at all?                             |
| Cost per hard pass      | Among completed tasks, what did each completion cost?     |
| Wall time per hard pass | Among completed tasks, how long did each completion take? |

**Table 2.** Co-primary endpoints.

Signal pass, “did the agent visibly use prior context”, sits next to these as the Chronicle-specific quality check. Equal hard pass with higher signal means Chronicle changed *what* gets done, not *whether* it gets done.

Repository-search per hard pass, the count of Read, Grep, and Glob tool calls per successful tuple, is the most Chronicle-specific cost-adjacent metric. It captures the “did memory point the agent at the right place before broad repository discovery” question that lives between latency and quality.

A *clean tuple* has empty `taint_kinds`. A *paired comparison* has matching (`track`, `scenario_id`, `seed`) across treatment and control, both clean. Bootstrap 95% CIs are over paired comparisons.

## 6 Evidence ladder

The three layers below sit on the same harness. The model is `claude-sonnet-4-5`. Isolation is Docker. Chronicle is served by a per-tuple sidecar daemon over TCP. Execution authentication is OAuth-only for the Claude Code subprocess, verified `apiKeySource`: `"none"` across all streams. Each layer answers a different shape of the same question, and each layer is stricter than the one before it.

## 6.1 Layer 1: vague, open-ended prior-context prompts

Vague tests open-ended user prompts mined from real sessions, where the answer the user wanted depends on prior-session decisions, failures, or integration details. Hard pass and signal pass are now scored by concept-group checks rather than the old “wrote a non-empty answer.md” floor, so completion and quality are now meaningfully separable.

The current reportable filter excludes four scenarios. Three are baseline-side cost or wall-time outliers where baseline timeouts and cost spikes would have inflated Chronicle’s relative advantage on the headline wall and cost numbers. The fourth is a documented weak-differential scenario where baseline can discover the key fact from current repository state. Including the four scenarios would widen Chronicle’s wall and cost margins, not narrow them; the exclusion is therefore conservative for Chronicle.

| Variant   | Clean tuples | Hard pass | Signal pass | Total cost | Cost / hard | Cost / signal | Signal / \$ | Mean wall | Mean agent |
|-----------|--------------|-----------|-------------|------------|-------------|---------------|-------------|-----------|------------|
| baseline  | 70           | 50        | 20          | \$25.601   | \$0.512     | \$1.280       | 0.781       | 139.0 s   | 128.4 s    |
| Chronicle | 72           | 68        | 53          | \$22.081   | \$0.325     | \$0.417       | 2.400       | 117.2 s   | 114.8 s    |

**Table 3.** Vague layer: per-arm summary on the 24-scenario reportable set.

| Metric                 | Pairs | Mean     | Median   | Bootstrap 95% CI     |
|------------------------|-------|----------|----------|----------------------|
| hard pass delta        | 70    | +0.229   | 0.000    | [+0.129, +0.329]     |
| signal pass delta      | 70    | +0.443   | 0.000    | [+0.329, +0.557]     |
| cost delta, all pairs  | 70    | −\$0.054 | −\$0.056 | [−\$0.117, +\$0.012] |
| wall delta, all pairs  | 70    | −21.2 s  | −11.4 s  | [−44.6 s, +2.6 s]    |
| agent delta, all pairs | 70    | −13.1 s  | −10.1 s  | [−36.5 s, +11.3 s]   |

**Table 4.** Vague layer: paired effects on 70 clean pairs.

Pass-conditioned deltas isolate the “when baseline succeeded” rows, where comparison is most fair to baseline:

| Condition              | Pairs | Mean cost $\Delta$ | Median cost $\Delta$ | Mean wall $\Delta$ | Median wall $\Delta$ |
|------------------------|-------|--------------------|----------------------|--------------------|----------------------|
| baseline hard-passed   | 50    | −\$0.064           | −\$0.047             | −21.0 s            | −8.3 s               |
| baseline signal-passed | 20    | −\$0.084           | −\$0.056             | −35.7 s            | −14.4 s              |

**Table 5.** Vague layer: pass-conditioned cost and wall deltas.

Three things stand out. First, hard pass moves from 50/70 (71%) to 68/72 (94%), and signal pass moves from 20/70 (29%) to 53/72 (74%). Second, the cost CI is consistent with a small win on average but not powered, while the conditioned cost delta is clearly negative when baseline succeeded: Chronicle was cheaper and faster on the seeds baseline could complete, in addition to completing 18 more. Third, there are no clean signal-loss cells in the dominance buckets across 70 paired comparisons. The vague layer earns a “more answers, no regressions” claim, with the pass-conditioned cost evidence as a sanity check that the gain is not a “Chronicle just spends more” artifact.

The caveat is that vague concept-check scoring is stronger than the old non-empty answer.md floor but not as strict as a blinded human grade, so public claims should still acknowledge manual quality work separately.

## 6.2 Layer 2: behaviour, project rules and conventions, including Headroom

Behaviour tests whether retrieval helps the agent follow project-specific rules and conventions mined from prior sessions. This is the layer where Chronicle is compared against a real competing memory

| Comparison            | Pairs | Win+cheaper | Win+costlier | Tie+cheaper | Tie+costlier | Loss+cheaper | Loss+costlier |
|-----------------------|-------|-------------|--------------|-------------|--------------|--------------|---------------|
| Chronicle vs baseline | 70    | 15          | 16           | 30          | 9            | 0            | 0             |

**Table 6.** Vague layer: dominance buckets, quality = signal pass.

tool, [Headroom](#), as well as stock Claude Code. Headroom is a meaningfully different design: offline learn --apply distills a curated CLAUDE.md plus MEMORY.md and a local HTTP proxy injects that summary into every model API call. Chronicle is per-prompt and lossy-but-fresh; Headroom is upfront and distilled-into-rules. Beating both is more informative than beating only one.

| Variant   | Clean tuples | Hard pass | Signal pass | Total cost | Cost / hard | Cost / signal | Signal / \$ | Mean agent |
|-----------|--------------|-----------|-------------|------------|-------------|---------------|-------------|------------|
| baseline  | 111          | 103       | 63          | \$7.208    | \$0.070     | \$0.114       | 8.741       | 38.1 s     |
| Headroom  | 112          | 101       | 74          | \$6.941    | \$0.069     | \$0.094       | 10.661      | 33.5 s     |
| Chronicle | 114          | 111       | 102         | \$7.609    | \$0.069     | \$0.075       | 13.405      | 32.4 s     |

**Table 7.** Behaviour layer: per-arm summary.

| Metric                 | Pairs | Mean     | Median   | Bootstrap 95% CI     |
|------------------------|-------|----------|----------|----------------------|
| hard pass delta        | 111   | +0.054   | 0.000    | [+0.018, +0.099]     |
| signal pass delta      | 111   | +0.333   | 0.000    | [+0.252, +0.423]     |
| cost delta, all pairs  | 111   | +\$0.003 | -\$0.003 | [-\$0.005, +\$0.011] |
| wall delta, all pairs  | 111   | -10.2 s  | -9.6 s   | [-16.4 s, -3.5 s]    |
| agent delta, all pairs | 111   | -5.2 s   | -6.7 s   | [-10.9 s, +1.1 s]    |

**Table 8.** Chronicle vs baseline (paired, clean) on the behaviour layer.

Hard-pass lift over baseline is +5.4 points, CI [+1.8, +9.9]; signal lift is +33.3 points, CI [+25.2, +42.3]. Hard-pass lift over Headroom is +7.1 points, CI [+2.7, +12.5]; signal lift is +24.1 points, CI [+17.0, +32.1]. Per-pair cost is approximately neutral against both arms (paired CIs cross zero). Wall time is meaningfully lower against baseline: -10.2 s, CI [-16.4, -3.5]. Against Headroom, agent time is the fair task-time alignment because the Headroom HTTP proxy adds wall-time overhead that is not attributable to the agent’s own work; on agent time, Chronicle and Headroom are effectively tied while Chronicle is ahead on hard pass and signal.

Behaviour is also where the dominance-bucket result is most legible. Across 223 clean paired comparisons against baseline and Headroom combined, there are zero loss cells. Chronicle did not produce a clean signal regression on the behaviour layer in this run.

### 6.3 Layer 3: chronicle\_advantage\_pack, three arms

The advantage pack is the layer designed end to end for the claim, run with three arms: baseline, chronicle-on, and headroom-on. Conditions match the previous layers: Sonnet model, Docker isolation, OAuth-only Claude Code subprocess, and the per-tuple sidecar daemon for Chronicle. Headroom uses its own offline learn --apply flow plus the local HTTP proxy. Scenarios are calibrated against ideal/ and noop/ fixtures so retrieval is what changes the outcome, and source cutoff prevents eval-authoring leakage.

Headroom wall time is intentionally omitted from the comparison tables because Headroom proxy and learn boot is not aligned with the baseline / Chronicle timing path. Agent time, the parsed duration of the Claude subprocess itself, is reported instead.

The triple-clean overlap covers 113 paired tuples where all three arms are present and clean. Two

| Metric                 | Pairs | Mean     | Median   | Bootstrap 95% CI     |
|------------------------|-------|----------|----------|----------------------|
| hard pass delta        | 112   | +0.071   | 0.000    | [+0.027, +0.125]     |
| signal pass delta      | 112   | +0.241   | 0.000    | [+0.170, +0.321]     |
| cost delta, all pairs  | 112   | +\$0.005 | -\$0.001 | [-\$0.006, +\$0.015] |
| agent delta, all pairs | 112   | -1.4 s   | -4.1 s   | [-7.6 s, +5.3 s]     |

**Table 9.** Chronicle vs Headroom (paired, clean) on the behaviour layer.

| Comparison            | Pairs | Win+cheaper | Win+costlier | Tie+cheaper | Tie+costlier | Loss+cheaper | Loss+costlier |
|-----------------------|-------|-------------|--------------|-------------|--------------|--------------|---------------|
| Chronicle vs baseline | 111   | 18          | 19           | 43          | 31           | 0            | 0             |
| Chronicle vs Headroom | 112   | 13          | 14           | 44          | 41           | 0            | 0             |

**Table 10.** Behaviour layer: dominance buckets, quality = signal pass.

Headroom tuples are dropped for `env_probe` taints in one behaviour scenario; no Chronicle or baseline tuples are dropped. Three Headroom tuples in the conversation subtrack additionally write native Claude memory files into the agent’s isolated home (7 writes total). Chronicle and baseline write zero native Claude memory across the entire pack.

| Variant   | Tuples | Hard          | Signal       | Cost      | Cost / hard | Cost / signal | Agent / hard | Search / hard |
|-----------|--------|---------------|--------------|-----------|-------------|---------------|--------------|---------------|
| baseline  | 113    | 79/113 (70%)  | 33/113 (29%) | \$12.2763 | \$0.1554    | \$0.3720      | 78.1 s       | 7.68          |
| Chronicle | 113    | 109/113 (96%) | 70/113 (62%) | \$13.9823 | \$0.1283    | \$0.1997      | 64.1 s       | 5.03          |
| Headroom  | 113    | 79/113 (70%)  | 36/113 (32%) | \$12.7253 | \$0.1611    | \$0.3535      | 90.7 s       | 8.16          |

**Table 11.** Advantage pack triple-clean overlap (113 paired tuples per arm).

Three things stand out. First, Headroom does not lift hard pass over baseline on the advantage pack: 79/113 in both arms. Second, on signal Headroom adds three passes over baseline (36 vs 33), where Chronicle adds 37 (70 vs 33). Third, Headroom is more expensive per hard pass (\$0.1611 vs baseline \$0.1554, vs Chronicle \$0.1283), uses more repository search per hard pass (8.16 vs baseline 7.68, vs Chronicle 5.03), and spends more agent time per hard pass (90.7 s vs baseline 78.1 s, vs Chronicle 64.1 s). On the advantage pack, Headroom helps marginally on signal and not at all on completion, while costing more and exploring more than baseline. This contrasts with the behaviour layer, where Headroom did lift signal over baseline. The behaviour layer is mostly front-of-turn `UserPromptSubmit` work; the advantage pack is conversation-weighted, and Headroom’s offline distilled `CLAUDE.md` plus `MEMORY.md` injection appears to lose more of its lift on the mid-task surface.

Chronicle’s per-success efficiency is now fully stated against both arms. Cost per hard pass is 17% below baseline and 20% below Headroom; cost per signal pass is 46% below baseline and 44% below Headroom; agent time per hard pass is 18% below baseline and 29% below Headroom; repository-search per hard pass is 35% below baseline and 38% below Headroom.

Chronicle versus baseline paired effects, computed on the 115-pair Chronicle / baseline subset (the two Headroom-tainted tuples are still paired here because they are clean for Chronicle and baseline):

On the headline endpoints, hard pass per pair and signal pass per pair, Chronicle versus baseline clears zero (+0.261 hard, +0.322 signal). On the same-task pairwise rows where both arms succeeded, Chronicle is slightly more expensive per pair (+\$0.0210, CI [+0.0106, +\$0.0320]) and slightly slower per pair (+8.9 s, CI [+4.3, +13.9]). Repository-search per pair on the same task is approximately neutral.

The two views are not in tension. Chronicle attempts and completes harder work per tuple, so it spends modestly more on each tuple while finishing many more tuples. Per-success spend is therefore lower

| Paired metric                         | Pairs | Mean      | Median    | Bootstrap 95% CI       |
|---------------------------------------|-------|-----------|-----------|------------------------|
| Full hard delta                       | 115   | +0.261    | 0.000     | [+0.183, +0.348]       |
| Full signal delta                     | 115   | +0.322    | 0.000     | [+0.235, +0.409]       |
| Cost delta, all clean pairs           | 115   | +\$0.0150 | +\$0.0085 | [\$0.0033, +\$0.0272]  |
| Cost delta, both hard-passed          | 81    | +\$0.0210 | +\$0.0120 | [\$0.0106, +\$0.0320]  |
| Wall delta, both hard-passed          | 81    | +8.9 s    | +3.9 s    | [+4.3 s, +13.9 s]      |
| Agent delta, both hard-passed         | 81    | +9.0 s    | +3.1 s    | [+4.3 s, +13.8 s]      |
| Repo-search delta, both hard-passed   | 81    | -0.111    | 0.000     | [-0.815, +0.556]       |
| Cost delta, both signal-passed        | 33    | +\$0.0028 | +\$0.0025 | [-\$0.0086, +\$0.0123] |
| Wall delta, both signal-passed        | 33    | +3.9 s    | +2.4 s    | [+1.5 s, +6.9 s]       |
| Repo-search delta, both signal-passed | 33    | +0.152    | 0.000     | [-0.455, +0.818]       |

**Table 12.** Advantage pack: Chronicle vs baseline paired effects.

even though per-pair spend on the same task is slightly higher. This is a “tasks solved per budget” result, not a “same task cheaper” result.

### Subtrack split, three arms

| Variant   | Hard  | Signal | Cost / hard | Agent / hard | Search / hard |
|-----------|-------|--------|-------------|--------------|---------------|
| baseline  | 19/22 | 9/22   | \$0.0725    | 37.2 s       | 4.53          |
| Chronicle | 22/22 | 21/22  | \$0.0686    | 32.2 s       | 3.36          |
| Headroom  | 19/22 | 15/22  | \$0.0829    | 44.6 s       | 4.42          |

**Table 13.** Behaviour subtrack of the advantage pack (22 paired tuples per arm). Two seeds of one behaviour scenario are dropped for env\_probe Headroom taints.

| Variant   | Hard  | Signal | Cost / hard | Agent / hard | Search / hard |
|-----------|-------|--------|-------------|--------------|---------------|
| baseline  | 60/91 | 24/91  | \$0.1816    | 91.1 s       | 8.68          |
| Chronicle | 87/91 | 49/91  | \$0.1434    | 72.2 s       | 5.45          |
| Headroom  | 60/91 | 21/91  | \$0.1858    | 105.4 s      | 9.35          |

**Table 14.** Conversation subtrack of the advantage pack (91 paired tuples per arm).

The conversation subtrack carries most of the per-success efficiency story. Chronicle’s cost per hard pass is 21% below baseline and 23% below Headroom; agent time per hard pass is 21% below baseline and 31% below Headroom; repository-search per hard pass drops 37% below baseline and 42% below Headroom. Headroom on the conversation surface is materially worse than baseline on signal (21 vs 24) and on every other reported metric, despite spending more agent time. The behaviour subtrack reverses the Headroom-versus-baseline picture: Headroom does lift signal over baseline on front-of-turn work (15 vs 9), but Chronicle still leads (21 vs 15) at lower cost and shorter agent time.

Chronicle versus baseline paired effects on the underlying current-table view (the per-pair hard, signal, cost, and wall CIs are unchanged from the two-arm reporting because those CIs do not depend on the Headroom arm): behaviour subtrack hard delta +0.125, CI [0.000, +0.292]; behaviour subtrack signal delta +0.500, CI [+0.292, +0.708]; conversation subtrack hard delta +0.297, CI [+0.209, +0.396]; conversation subtrack signal delta +0.275, CI [+0.187, +0.363]. Both subtracks are necessary: behaviour shows the claim survives the front-of-turn surface, and conversation shows it survives the harder mid-task surface where injection fires repeatedly.



## 7 **Headline figure: budget dominance**

The natural visualisation for a “tasks solved per budget” claim is a budget-dominance curve. The  $x$ -axis is a cost or wall-time budget; the  $y$ -axis is the count of tasks solved at or below that budget.

The shape this data supports is: across the practical range of per-task cost budgets and per-task wall-time budgets, Chronicle’s solved count dominates baseline’s solved count at every budget. There is no budget at which baseline solves more tasks than Chronicle on this pack. The dominance gap widens as the budget grows, because Chronicle continues to convert harder seeds while baseline plateaus at 79/113 hard pass and 33/113 signal pass on the triple-clean overlap.

This figure replaces the previous “mean cost delta” framing because mean cost delta on the same pair systematically underweights tasks the baseline could not complete, which is exactly where Chronicle’s advantage lives. Budget dominance is the right shape for the claim.

## 8 **Cross-run claim ladder**

The three layers support three layers of claim, in increasing strictness:

1. **Vague, current reportable.** Chronicle answers open-ended prior-context prompts with much higher hard- and signal-pass rates, lower total cost, lower cost per hard pass, lower cost per signal pass, no clean signal-loss cells, and a clearly negative pass-conditioned cost and wall delta when baseline succeeded.
2. **Behaviour, current reportable.** Chronicle beats both baseline and Headroom on hard pass, signal pass, and signal per dollar. Wall time is meaningfully lower against baseline; agent time is the fair Headroom alignment. Per-pair cost is neutral. Zero clean signal-loss cells across 223 paired comparisons.
3. **Advantage pack with Headroom.** Chronicle solves materially more tasks than both stock Claude Code and Headroom on the purpose-built three-arm pack, with lower cost per solved task and lower agent time per solved task against both arms. Headroom does not lift hard pass over baseline on the pack and adds only a marginal signal lift, while costing more and exploring more. Three Headroom tuples in the conversation subtrack write native Claude memory files; Chronicle and baseline write zero. The honest framing is “tasks solved per budget”, supported by the budget-dominance curve, with same-task pairwise rows showing modestly positive cost and wall-time deltas that are consistent with Chronicle attempting and completing harder work per tuple.

Each layer is a stricter test than the one before. Vague tests open-ended prior-context prompts with the easiest scoring. Behaviour tests rule and convention following with a real competing memory tool in the room. The advantage pack tests a purpose-built scenario set with conversation-time injection and a strict hard / signal split.

## 9 **What this earns and what it does not**

The three layers earn the defensible claim above. They do not yet earn the bolder framing of “Chronicle is faster and cheaper on every coding task.” When baseline already succeeds, Chronicle’s same-task per-pair cost is approximately neutral. Chronicle does not make easy successful tasks cheaper; it makes hard tasks possible. A stronger general-purpose coding agent on a self-contained refactor should already do well, and expecting Chronicle to beat it there is the wrong test.

Specific limits to keep visible:

- **Per-pair same-task cost is positive, not negative.** On the advantage-pack same-task rows, Chronicle

is +\$0.0210 mean and +8.9 s mean, with CIs that do not cross zero. The Chronicle win lives in solving more tasks at comparable per-pair spend, not in making the same task cheaper.

- **Vague concept-check scoring is stronger than the old answer .md floor but not as strict as blinded human grading.** Manual quality work should be cited separately alongside any public claim.
- **Headroom is fairly set up but not deeply tuned.** Headroom was run with its canonical `learn --apply flow` and the local HTTP proxy, on the same Docker isolation, OAuth-only authentication, and Sonnet model as the other arms. Its conversation-surface underperformance against baseline is plausibly a real artefact of design rather than a setup issue: distilled `CLAUDE.md` and `MEMORY.md` injected on every API call can derail multi-turn exploration when any retrieved fact is tangential to the current sub-task. Behaviour-layer evidence, where Headroom does lift signal over baseline, suggests the substrate is functioning rather than starved.
- **Single model.** All three layers run on `claude-sonnet-4-5`. Behaviour on Opus or Haiku is untested.

## 10 Limitations

Chronicle is operationally usable but not yet polished. Implementation gaps include incremental JSONL refresh on the `compose.enabled` path, scoring writeback for `USER_AUTHORED_BOOST` and citation priors, hard `PreToolUse` blocking for pinned policy, and broader action-mining coverage. Architecturally, dense retrieval can still mix up negation, time, and replaced claims; pre-turn retrieval can miss because the model has not reasoned about the task yet; several keyword lists assume English and common HTTP and POSIX wording; and large-project cost is dominated by typed-claim LLM extraction rather than local chunk indexing. Eval-harness gaps include single-model and single-seed-count coverage on the advantage pack and taint-scanner blind spots on host-transcript leak paths. The package itself is currently a research monorepo rather than a clean public release.

## 11 Further work

**Multi-model matrix.** All current numbers are on `claude-sonnet-4-5`. A natural follow-up runs the same pack on Opus and Haiku to test whether the per-success efficiency story holds across model scales. The hypothesis is that the headline shape is preserved but the absolute numbers compress on Opus (where the baseline is already strong) and dilate on Haiku (where prior context is more useful).

**Petri-style memory audits.** Anthropic’s [Petri](#) framework is an alignment-auditing tool that deploys an automated auditor agent to test a target model through multi-turn conversations with simulated users and tools, then has a judge score the resulting transcripts across multiple dimensions. Petri itself is not a coding-memory benchmark, but the harness shape is the relevant pointer. A Chronicle version would seed each scenario with the memory challenge under test (stale memories, missing prior decisions, misleading near-matches, retrieved chunks the agent should ignore), and a Petri-style auditor would drive the multi-turn coding interaction. The judge would score not just final correctness but whether the agent relied on the right memory, ignored bad memory, or wasted turns rediscovering context. That complements scripted check harnesses rather than replacing them.

**Citation faithfulness.** Recent RAG citation work distinguishes citation correctness from citation faithfulness: a cited source can support a claim even if the model did not rely on it. Chronicle can join injection logs, citation logs, file edits, and Bash calls to ask, when memory was injected and cited, whether the agent’s action actually depended on it. If unfaithful citations predict failed checks, citation faithfulness becomes a useful signal before a run finishes rather than an after-the-fact paper metric.